

# RESSL UVM Sequences to the Mat

Jeff McNeal, Bryan Morris

Verilab

Jeff McNeal (Verilab US), Bryan Morris (Verilab Canada)

[www.verilab.com](http://www.verilab.com)

## ABSTRACT

*Read-Evaluate-Start-Sequence-Loop (RESSL -- pronounced “wrestle”) is inspired by the Read-Evaluate-Print-Loop (REPL) found in Lisp and Python. The REPL in these languages encourage a rapid, iterative and interactive development process allowing the user to easily develop and test new sequences with a minimum of overhead.*

*In the context of ASIC verification, RESSL enables the iterative development and debug of UVM sequences. Similar to the Lisp REPL, it includes four phases:*

***Read:*** A simple interpreter allowing the user to input commands via STDIN.

***Evaluate:*** The evaluator takes those commands and executes them. These commands include among others, the ability to clone, alter parameters and start sequences.

***Start-Sequence:*** The system starts the sequence (and any sub-sequences) defined.

***Loop:*** Clean up and return back to the Read.

*This paper provides details on the usage model, implementation and future work planned for the RESSL.*

## Table of Contents

Introduction.....	4
Sequences are great, but.....	4
The Lisp Inspiration.....	5
RESSL Usage.....	6
Example Usage Scenario .....	7
RESSL API .....	15
Typical Usage Scenarios .....	16
RESSL Design .....	16
Interpreter.....	17
Get User Input.....	17
Tokenizer .....	17
Verbs .....	18
Start Sequence.....	18
Introspection/Reflection.....	19
Storage And Retrieval.....	19
Interesting... but how hard is it to add to my existing environment?.....	19
Step 1. Add the svlib .....	20
Step 2. Declare and instantiate the RESSL object .....	20
Step 3. Disable default sequence .....	20
Step 4. Seed RESSL's sequence registry with the lrmw_seq .....	21
Step 5. Start RESSL.....	21
UVM Library Modifications.....	21
Modifying uvm_object Class.....	21
Modifying UVM Field Macros .....	21
Limitations and Future Work.....	23
We do not have a way to add new constraints to a sequence. ....	23
We do not provide a way to add new fields to sequences. ....	23
Future Features /Improvements .....	24
Conclusions.....	24
Acknowledgements.....	25
References.....	25

## **Table of Figures**

Figure 1 - RESSL Processing .....	6
Figure 2 - RESSL Class (UML Class Diagram) .....	17
Figure 3 - RESSL Command (UML Classs Diagram) .....	18
Figure 4 - RESSL External File .....	19

## **Table of Tables**

Table 1 - Current RESSL Commands.....	15
---------------------------------------	----

## Introduction

Dynamic languages such as Python [1], Ruby [2] and Lisp [3] provide ways to develop programs iteratively and interactively. Using a command line interpreter, programmers can create small snippets of functions that they can immediately test for correctness. They can then iteratively build and test these snippets until they have fully defined functions. This system is referred to in Lisp and Python as the REPL (Read-Evaluate-Print-Loop). Unfortunately, providing this kind of functionality is not possible in SystemVerilog as its compile and run model is vastly different from the dynamic languages referred to above.

However, UVM can generate new objects using a factory pattern [4], and UVM sequences can nest sub-sequences within sequences -- thus enabling dynamic creation of progressively more complicated sequences. These sequences can be started dynamically on a compatibly-typed sequencer. Combining the REPL development methodology with UVM's inherent ability to create new sequences using a factory pattern and launch sequences dynamically, we created a solution we call RESSL (Read-Evaluate-Start Sequence-Loop).

RESSL adds a simple parser, interpreter and execution loop that allows the user to interactively create new sequences (from sequences that are already declared, compiled and registered with the `uvm_factory`), build new composite sequences from other sequences, and then run these sequences -- all done while the simulation is running.

In addition to providing the means to create and run sequences interactively, we made updates to many of the UVM field macros (`uvm_field_*`) to allow the user to view a sequence's variables (which have been defined using the field macro) and update its value. While this introspection/reflection capability is quite useful in this context, it is a key limitation for those groups who are not able to patch their UVM library. However, it is possible to add this capability without modifying UVM -- it simply requires more coding than the transparent 'under-the-hood' solution of modifying the UVM field macros.

## Sequences are great, but...

UVM Sequences are a powerful feature of UVM. Using the UVM sequence classes and associated macros, coupled with SystemVerilog random constraints, sequences of interesting scenarios can be created quickly and easily. More importantly, complicated sequences can easily and iteratively be created based on simpler sequences that already exist.

One limitation with this capability is that new sequences must be coded and then compiled before they can be used. Since they are compiled source code, they cannot be modified or expanded while a simulation is running.

There are two principal users of sequences: RTL design engineers and Verification engineers and they each have different use cases. RTL designers want sequences that are easy to create and modify while testing their code. They generally don't want to mess with the testbench or learn UVM in order to do simple testing. Ideally they would like to be able to quickly build and tweak stimulus, in this case a sequence, on their own without waiting for the verification engineer to modify the sequence source code or create a multitude of command line arguments to set different attributes of the sequence. On the other hand, verification engineers want it to be

convenient to be able to change the attributes of the sequence under development without having to recompile the code and then launch the sequence to verify its behaviour. For both use cases, it would be ideal that once you have created an interesting sequence you should be able to save and replay it for subsequent simulations or regressions.

The system we have created, RESSL, provides the ability to dynamically create sequences from existing sequences, tweak any sequence's parameters, and then start the interactively created sequences, all while the simulation is running from a command line prompt using a simple API.

## The Lisp Inspiration

One of the original languages to provide dynamic, interactive development capability was Lisp. It was one of the first languages to introduce a REPL [8]. The four phases of a Lisp REPL are:

- **Read:** accept input from the user (typically Lisp code.)
- **Evaluate:** process the input from the user, ensure it is syntactically correct and meaningful in the current context and then allow it to be executed within the Lisp system.
- **Print:** the results from the Evaluate phase are then presented, allowing the developer to immediately and interactively build their program.
- **Loop:** returns back to the Read phase enabling the iterative process.

This type of development model would be useful when writing or using a SystemVerilog testbench. However, since SystemVerilog is a compiled, static language, we cannot interactively execute new code during run time. RESSL takes its inspiration from interactive languages and goes as far as it can to provide a way to interactively create, alter, inspect, and start UVM sequences. This provides much of the benefit of a full REPL in an interactive language in SystemVerilog. As the name implies, and similar to above, it comprises four phases as shown in Figure 1 - RESSL Processing on page 6:

- **Read:** a simple interpreter that reads the command line input, tokenizes it and passes it to the Evaluate phase
- **Evaluate:** interprets the command and executes.  
This evaluation may include:
  - adding, changing the sequence registry<sup>1</sup> where all sequences are stored.
  - displaying, modifying or randomizing sequence variable values via introspection [5][6]
  - creating new sequences.
  - storing and loading sequences to and from files.
- **Start-Sequence:** starts the sequence specified from the sequence registry. Note that sequences can contain sub-sequences.
- **Loop:** executes any clean-up and loops back to the Read phase.

---

<sup>1</sup> The sequence registry is a structure to hold the sequences known to RESSL. More details are provided in the RESSL design on page 12.

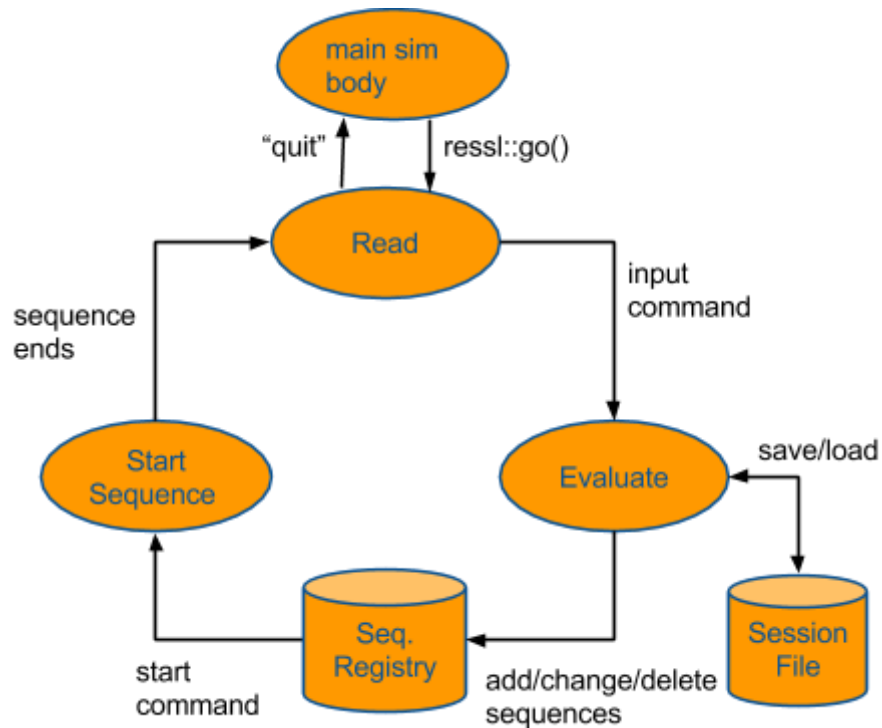


Figure 1 - RESSL Processing

## RESSL Usage

General UVM guidelines recommend that sequences be divided into several levels. At the bottom level are atomic or protocol specific sequences. These sequences generally perform a single, well bounded task, such as transferring a word of data, or setting control signals. At the next higher level are sequences that perform a task using the lower layer sequences. These may be register writes and reads, or packet transfers. The third level of sequences use the second (and first) levels to accomplish tasks, such as transfer a block of memory, or configure the DUT.

RESSL requires that at least the lowest level of sequences is already written and working in the verification environment. At that point the user can start up RESSL and interactively build mid-level sequences out of the low level sub-sequences in an interactive manner. RESSL allows the user to string together several sub-sequences in a sequence and execute the sequence, then observe the result.

Then changes can be made to the sequence and it can be run again, all within the same simulation execution. The user can change the values of fields in the sequence to whatever is desired, allowing users to quickly try out different combinations of values without editing the source or re-compiling.

Once the user is happy with the sequences that have been built, they can be saved to a file for later use, or to be coded into the verification environment as tests. In a future simulation the user can load the sequences and use them again.

## Example Usage Scenario

This scenario uses the ubus example provided in the UVM1.2 library. We have integrated RESSL into the testbench using the instructions provided in the section entitled “Interesting... but how hard is it to add to my existing environment?” on page 19. The environment includes several sequences to do a single write (`write_byte_seq`), a single read (`read_byte_seq`), and another sequence that creates a read-modify-write sequence (`lrmw_seq`).

Here we start up the simulation and once it runs the `ressl::go()` task it transfers control to the RESSL command line prompt which displays to the Unix stdout:

```
./simv +UVM_VERBOSITY=UVM_LOW -l vcs.log +UVM_TESTNAME=test_2m_4s
Chronologic VCS simulator copyright 1991-2013
Contains Synopsys proprietary information.
Compiler version H-2013.06-SP1_Full64; Runtime version H-2013.06-
SP1_Full64; May 11 19:29 2015
UVM_INFO ../../../../src/base/uvm_root.svh(392) @ 0: reporter
[UVM/RELNOTES]
-----
UVM-1.2
(C) 2007-2014 Mentor Graphics Corporation
(C) 2007-2014 Cadence Design Systems, Inc.
(C) 2006-2014 Synopsys, Inc.
(C) 2011-2013 Cypress Semiconductor Corp.
(C) 2013-2014 NVIDIA Corporation
-----

***** IMPORTANT RELEASE NOTES *****

You are using a version of the UVM library that has been compiled
with `UVM_NO_DEPRECATED undefined.
See http://www.eda.org/svdb/view.php?id=3313 for more details.

You are using a version of the UVM library that has been compiled
with `UVM_OBJECT_DO_NOT_NEED_CONSTRUCTOR undefined.
See http://www.eda.org/svdb/view.php?id=3770 for more details.

(Specify +UVM_NO_RELNOTES to turn off this notice)

UVM_INFO @ 0: reporter [RNTST] Running test test_2m_4s...
UVM_INFO test_lib.sv(67) @ 0: uvm_test_top [test_2m_4s] Printing the test
topology :
-----
Name                                     Type                                     Size  Value
-----
uvm_test_top                             test_2m_4s                             -      @350
  ubus_example_tb0                       ubus_example_tb                        -      @388
    scoreboard0                          ubus_example_scoreboard                -      @425
      item_collected_export              uvm_analysis_imp                       -      @434
      disable_scoreboard                  integral                               1       'h0
      num_writes                          integral                               32      'd0
      num_init_reads                      integral                               32      'd0
      num_uninit_reads                    integral                               32      'd0
      recording_detail                    uvm_verbosity                          32      UVM_FULL
    ubus0                                 ubus_env                               -      @416
```

bus_monitor	ubus_bus_monitor	-	@449
masters[0]	ubus_master_agent	-	@484
masters[1]	ubus_master_agent	-	@497
slaves[0]	ubus_slave_agent	-	@509
slaves[1]	ubus_slave_agent	-	@518
slaves[2]	ubus_slave_agent	-	@527
slaves[3]	ubus_slave_agent	-	@536
has_bus_monitor	integral	1	'h1
num_masters	integral	32	'h2
num_slaves	integral	32	'h4
intf_checks_enable	integral	1	'h1
intf_coverage_enable	integral	1	'h1
recording_detail	uvm_verbosity	32	UVM_FULLL
recording_detail	uvm_verbosity	32	UVM_FULLL

---

NOTE: We populate our sequence and sequencer repository with an existing sequence and sequencer from within the testbench itself.

```
seq loop_read_modify_write_seq (type=loop_read_modify_write_seq) added.
seqr sequencer (type=uvm_sequencer) added.
```

---

Waiting for user input.

```
SV-RESSL: Pinning Sequences to the Mat since 2014...
```

```
[*] >>>
```

---

Simple help

```
[*] >>> help
```

```
--- RESSL Help ---
```

```
add <sequence name> [repeat_count]
attach <seqr> <to_seq> [ALL | <seq_index>]
create <SEQ | SEQR> <type> <name>
describe <sequence name>
help [verb]
list
load <filename>
move <from_index> <to_index>
quit = exits RESSL and continues the simulation.
randomize <seq_index>
save <filename>
select <seq_path>
set <seq_index> <field> <value>
start [repeat_count] = Execute the currently selected sequence
[repeat_count] times.
```

---

Let's list the known UVM sequences and sequencers currently defined in the registry.

```
[*] >>> list
Sequence Registry:
    lrmw_seq [#subseq:1]
Sequencer Registry:
```



```
ubus_example_tb0.ubus0.masters[0].sequencer (type:uvm_sequencer)
```

---

So far in this example, we have one sequence named `lrmw_seq` and one sequencer which have been added into the sequence and sequencer repository in the testbench's SystemVerilog code. These are now part of the sequence and sequencer registry holding all the UVM sequences and sequencers known to RESSL.

Create a new UVM sequence of type `write_byte_seq` called `wbs`, and see that it is automatically added to the sequence repository.

---

```
[*] >>> create SEQ write_byte_seq wbs
seq wbs (type=wbs) added.

[*] >>> list
Sequence Registry:
    lrmw_seq [#subseq:1]
    wbs [#subseq:1]
Sequencer Registry:
    ubus_example_tb0.ubus0.masters[0].sequencer (type:uvm_sequencer)
```

---

Select one of the existing sequences to work with:

```
[*] >>> select lrmw_seq
```

---

Using RESSL's introspection capabilities, describe the current contents of the fields of the sequence (specifically, the ones defined using `uvm_field_*` macros).

```
[lrmw_seq] >>> describe
[0] Sequence: lrmw_seq (type:loop_read_modify_write_seq) [Sequencer type:
unattached]
    Fields:
        Field: itr = 0
```

As can be seen above, the `lrmw_seq` has only one existing sequence (also called `lrmw_seq`), with only one field called `itr` that is current set to 0. Set the field to new value:

```
[lrmw_seq] >>> set 0 itr 62

[lrmw_seq] >>> describe
[0] Sequence: lrmw_seq (type:loop_read_modify_write_seq) [Sequencer type:
unattached]
    Fields:
        Field: itr = 62
```

---

Now, let's select the previously created write sequence and start it on the default UVM sequencer.

```
[lrmw_seq] >>> select wbs
```

```
[wbs] >>> describe
```

```
[0] Sequence: wbs (type:write_byte_seq)[Sequencer type: uvm_sequencer]
      Fields: Empty
```

```
[wbs] >>> start
```

Name	Type	Size	Value
sequencer	uvm_sequencer	-	@573
rsp_export	uvm_analysis_export	-	@582
recording_detail	uvm_verbosity	32	UVM_FULL
seq_item_export	uvm_seq_item_pull_imp	-	@700
recording_detail	uvm_verbosity	32	UVM_FULL
recording_detail	uvm_verbosity	32	UVM_FULL
arbitration_queue	array	0	-
lock_queue	array	0	-
num_last_reqs	integral	32	'd1
num_last_rsps	integral	32	'd1

```
UVM_INFO ../sv/ubus_bus_monitor.sv(223) @ 3480:
```

```
uvm_test_top.ubus_example_tb0.ubus0.bus_monitor [ubus_bus_monitor] Transfer
collected :
```

Name	Type	Size	Value
ubus_transfer_inst	ubus_transfer	-	@459
addr	integral	16	'h0
read_write	ubus_read_write_enum	32	WRITE
size	integral	32	'h1
data	da(integral)	1	-
[0]	integral	8	'h0
wait_state	da(integral)	1	-
error_pos	integral	32	'h0
transmit_delay	integral	32	'h0
master	string	10	masters[0]
slave	string	9	slaves[0]
begin_time	time	64	3440
end_time	time	64	3480

```
UVM_INFO ubus_example_scoreboard.sv(89) @ 3480:
```

```
uvm_test_top.ubus_example_tb0.scoreboard0 [ubus_example_scoreboard] WRITE to
existing address...Updating address : 0 with data : 0
```

Create another sequence called rbs (of type read\_byte\_seq), and start it on the default sequencer.

```
[wbs] >>> create SEQ read_byte_seq rbs
seq rbs (type=rbs) added.
```

```
[wbs] >>> select rbs
```

```
[rbs] >>> start
```

Name	Type	Size	Value
------	------	------	-------

```

sequencer          uvm_sequencer          -      @573
  rsp_export        uvm_analysis_export    -      @582
    recording_detail uvm_verbosity          32     UVM_FULL
  seq_item_export    uvm_seq_item_pull_imp    -      @700
    recording_detail uvm_verbosity          32     UVM_FULL
  recording_detail    uvm_verbosity          32     UVM_FULL
  arbitration_queue  array                0      -
  lock_queue         array                0      -
  num_last_reqs      integral              32     'd1
  num_last_rsps      integral              32     'd1

```

```

UVM_INFO ../sv/ubus_bus_monitor.sv(223) @ 3550:
uvm_test_top.ubus_example_tb0.ubus0.bus_monitor [ubus_bus_monitor] Transfer
collected :

```

Name	Type	Size	Value
ubus_transfer_inst	ubus_transfer	-	@459
addr	integral	16	'h0
read_write	ubus_read_write_enum	32	READ
size	integral	32	'h1
data	da(integral)	1	-
[0]	integral	8	'h0
wait_state	da(integral)	1	-
error_pos	integral	32	'h0
transmit_delay	integral	32	'h0
master	string	10	masters[0]
slave	string	9	slaves[0]
begin_time	time	64	3510
end_time	time	64	3550

```

UVM_INFO ubus_example_scoreboard.sv(75) @ 3550:
uvm_test_top.ubus_example_tb0.scoreboard0 [ubus_example_scoreboard] READ to
existing address...Checking address : 0 with data : 0

```

---

Create a NEW sequence that will be an aggregate of these two simpler sequences i.e., this sequence will do a write, followed by a read, and then another write.

First, create the write sequence called `wr_rd_wr_seq`, and then add the existing read and write sequences.

```

[rbs] >>> create SEQ write_byte_seq wr_rd_wr_seq
seq wr_rd_wr_seq (type=wr_rd_wr_seq) added.

[rbs] >>> select wr_rd_wr_seq

[wr_rd_wr_seq] >>> add rbs

[wr_rd_wr_seq] >>> add wbs

```

---

Describe the new sequence to see the write (indexed by [0]) followed by a read (index = [1]), and then another write (index = [2]).

```

[wr_rd_wr_seq] >>> describe

```

```

[0] Sequence: wr_rd_wr_seq (type:write_byte_seq)[Sequencer type:
uvm_sequencer]
    Fields: Empty
[1] Sequence: wr_rd_wr_seq (type:read_byte_seq)[Sequencer type:
uvm_sequencer]
    Fields: Empty
[2] Sequence: wr_rd_wr_seq (type:write_byte_seq)[Sequencer type:
uvm_sequencer]
    Fields: Empty

```

---

Start the sequence and see the write...

```
[wr_rd_wr_seq] >>> start
```

Name	Type	Size	Value
sequencer	uvm_sequencer	-	@573
rsp_export	uvm_analysis_export	-	@582
recording_detail	uvm_verbosity	32	UVM_FULL
seq_item_export	uvm_seq_item_pull_imp	-	@700
recording_detail	uvm_verbosity	32	UVM_FULL
recording_detail	uvm_verbosity	32	UVM_FULL
arbitration_queue	array	0	-
lock_queue	array	0	-
num_last_reqs	integral	32	'd1
num_last_rsps	integral	32	'd1

```

UVM_INFO ../sv/ubus_bus_monitor.sv(223) @ 3620:
uvm_test_top.ubus_example_tb0.ubus0.bus_monitor [ubus_bus_monitor] Transfer
collected :

```

Name	Type	Size	Value
ubus_transfer_inst	ubus_transfer	-	@459
addr	integral	16	'h0
read_write	ubus_read_write_enum	32	WRITE
size	integral	32	'h1
data	da(integral)	1	-
[0]	integral	8	'h0
wait_state	da(integral)	1	-
error_pos	integral	32	'h0
transmit_delay	integral	32	'h0
master	string	10	masters[0]
slave	string	9	slaves[0]
begin_time	time	64	3580
end_time	time	64	3620

```

UVM_INFO ubus_example_scoreboard.sv(89) @ 3620:
uvm_test_top.ubus_example_tb0.scoreboard0 [ubus_example_scoreboard] WRITE to
existing address...Updating address : 0 with data : 0

```

... followed by a read...

Name	Type	Size	Value
------	------	------	-------

```

-----
sequencer          uvm_sequencer          -      @573
  rsp_export        uvm_analysis_export      -      @582
    recording_detail uvm_verbosity          32     UVM_FULL
  seq_item_export    uvm_seq_item_pull_imp      -      @700
    recording_detail uvm_verbosity          32     UVM_FULL
  recording_detail    uvm_verbosity          32     UVM_FULL
  arbitration_queue  array                      0      -
  lock_queue         array                      0      -
  num_last_reqs      integral                 32     'd1
  num_last_rsps      integral                 32     'd1
-----

```

UVM\_INFO ../sv/ubus\_bus\_monitor.sv(223) @ 3690:  
 uvm\_test\_top.ubus\_example\_tb0.ubus0.bus\_monitor [ubus\_bus\_monitor] Transfer collected :

```

-----
Name                Type                Size  Value
-----
ubus_transfer_inst  ubus_transfer          -      @459
  addr              integral                 16     'h0
  read_write        ubus_read_write_enum   32     READ
  size              integral                 32     'h1
  data              da(integral)             1      -
    [0]             integral                 8      'h0
  wait_state        da(integral)             1      -
  error_pos         integral                 32     'h0
  transmit_delay    integral                 32     'h0
  master            string                  10     masters[0]
  slave             string                  9      slaves[0]
  begin_time        time                   64     3650
  end_time          time                   64     3690
-----

```

UVM\_INFO ubus\_example\_scoreboard.sv(75) @ 3690:  
 uvm\_test\_top.ubus\_example\_tb0.scoreboard0 [ubus\_example\_scoreboard] READ to existing address...Checking address : 0 with data : 0

... followed by the last write...

```

-----
Name                Type                Size  Value
-----
sequencer          uvm_sequencer          -      @573
  rsp_export        uvm_analysis_export      -      @582
    recording_detail uvm_verbosity          32     UVM_FULL
  seq_item_export    uvm_seq_item_pull_imp      -      @700
    recording_detail uvm_verbosity          32     UVM_FULL
  recording_detail    uvm_verbosity          32     UVM_FULL
  arbitration_queue  array                      0      -
  lock_queue         array                      0      -
  num_last_reqs      integral                 32     'd1
  num_last_rsps      integral                 32     'd1
-----

```

UVM\_INFO ../sv/ubus\_bus\_monitor.sv(223) @ 3760:  
 uvm\_test\_top.ubus\_example\_tb0.ubus0.bus\_monitor [ubus\_bus\_monitor] Transfer collected :

---

Name	Type	Size	Value
ubus_transfer_inst	ubus_transfer	-	@459
addr	integral	16	'h0
read_write	ubus_read_write_enum	32	WRITE
size	integral	32	'h1
data	da(integral)	1	-
[0]	integral	8	'h0
wait_state	da(integral)	1	-
error_pos	integral	32	'h0
transmit_delay	integral	32	'h0
master	string	10	masters[0]
slave	string	9	slaves[0]
begin_time	time	64	3720
end_time	time	64	3760

---

```
UVM_INFO ubus_example_scoreboard.sv(89) @ 3760:
uvm_test_top.ubus_example_tb0.scoreboard0 [ubus_example_scoreboard] WRITE to
existing address...Updating address : 0 with data : 0
```

---

Save the operations from the current session. Save to the file `ressl.ex1` that can be reloaded in subsequent sessions.

```
[wr_rd_wr_seq] >>> save ressl.ex1
Saving[0]: create SEQ write_byte_seq wbs
Saving[1]: select wbs
Saving[2]: select lrmw_seq
Saving[7]: set 0 itr 62
Saving[8]: select wbs
Saving[9]: create SEQ read_byte_seq rbs
Saving[10]: select rbs
Saving[11]: create SEQ write_byte_seq wr_rd_wr_seq
Saving[12]: select wr_rd_wr_seq
Saving[13]: add rbs
Saving[14]: add wbs
Got it. Saved session to file
ressl.ex1... and cleared the command history. It's fresh and clean once
again.
```

---

Exit from RESSL, the test continues on from where RESSL was started.

```
[wr_rd_wr_seq] >>> quit
UVM_INFO ../../../../src/base/uvm_objection.svh(1271) @ 8760: reporter
[TEST_DONE] 'run' phase is ready to proceed to the 'extract' phase
```

---

## RESSL API

The current set of commands available in RESSL are:

**Table 1 - Current RESSL Commands**

Verb	Description
load <registry_file>	Load sequence registry from file.
store <registry_file>	Store sequence registry to a file.
select <seq_name>	Start recording a new sequence (or open an existing sequence); which is added to the sequence registry
List	List all the sequence names in the sequence registry
create <seq_type> <seq_name>	Using the uvm_factory to create a uvm_sequence_base of type seq_type and assigning it the name seq_name
add <seq_name> <repeat>	Adds the sub-sequence named seq_name to the current sequence. Optionally added repeat times.
copy <seq_name> <copy_name>	Copies the sequence to a new name: any sub-sequences are also copied. <i>Not currently available.</i>
delete [<seq_name>] [<seq_index>]	Delete the sub-sequence indexed by seq_index from the sequence named seq_name. If no seq_index is supplied, it deletes the entire seq_name from the sequence registry.
move <src_index> <dst_index>	Allows you to move a sub-sequence up or down to a new index position in the sequence.
describe <seq_name> [<index>]	Displays the attributes for named sequence's sub-sequence at the index specified. The attributes for all sub-sequences will be provided if the index is not supplied.
set <index> "field" "value"	Sets a field of one of the sub-sequences (at index) in the sequence.
start [seq_name] [index]	Starts the sequence. With no arguments it starts the currently selected sequence; Supplying only the seq_name starts that sequence. Supplying both arguments starts the sub-sequence indicated by index of the seq_name sequence.
shuffle <seq_name> [<shuffled_seq_name>]	In addition to starting the named sequence, it first randomizes the order of the sub-sequences defined to create a different sequence. By default, this is a transient shuffle where the original order defined remains--unless the user supplies the optional shuffled_seq_name in which case the shuffled sequence is copied to the new name.
help [cmd]+	Displays help message. With no arguments it displays all the known commands; otherwise it displays the help message for the <cmd> specified.
randomize <seq_index>	Randomize the sub-sequence at seq_index
attach <seqr> <seq>	Attach the specified sequencer seqr to the seq. This enables RESSL to access multiple sequences and sequencers within the same simulation.

## Typical Usage Scenarios

We envision RESSL being useful to a variety of people in a variety of ways.

- Simple sequences for RTL debug
- Debugging sequences, drivers, monitors.
- Developing and identifying an interesting set of sequences
- Quick way to develop a sequence library.

## RESSL Design

The design is conceptually split into four parts:

- *Interpreter*: The ‘R’ and ‘E’ of **RESSL**. A simple command execution loop that accepts input from the command line, tokenizes the input, and then creates and executes the appropriate command.
- *Sequence Loop*: The “SSL” of **RESSL**. Starts the sequences identified and loops back to the interpreter.
- *Introspection/Reflection*: The silent “I” in **RESSL**. A necessary part of the design is a crude introspection capability to enable dynamic setting of a sequence’s attributes via a command line.
- *Storage*: The ability to record, save and re-load libraries of previously developed sequences.

From the UML class diagram below there are currently two main classes that implement the RESSL functionality. The `ressl` class derives from `uvm_object` and provides the following functionality:

- **a sequence registry**: a simple database of all the defined `uvm_sequence_base` objects. Sequences are added into the registry using the `add_to_sequence_registry()` function or using the `add` command. The current implementation uses a simple associative array of handles to `uvm_sequence` objects keyed by the sequence name.
- **a sequencer registry**: similar to the sequence registry, except it holds handles to all the registered `uvm_sequencer` objects.
- **an interpreter** that accepts input from a simple command line display, tokenizes the input and then passes processing onto the associated `ressl_cmd` that implements the requested command.



## RESSL - Class Hierarchy

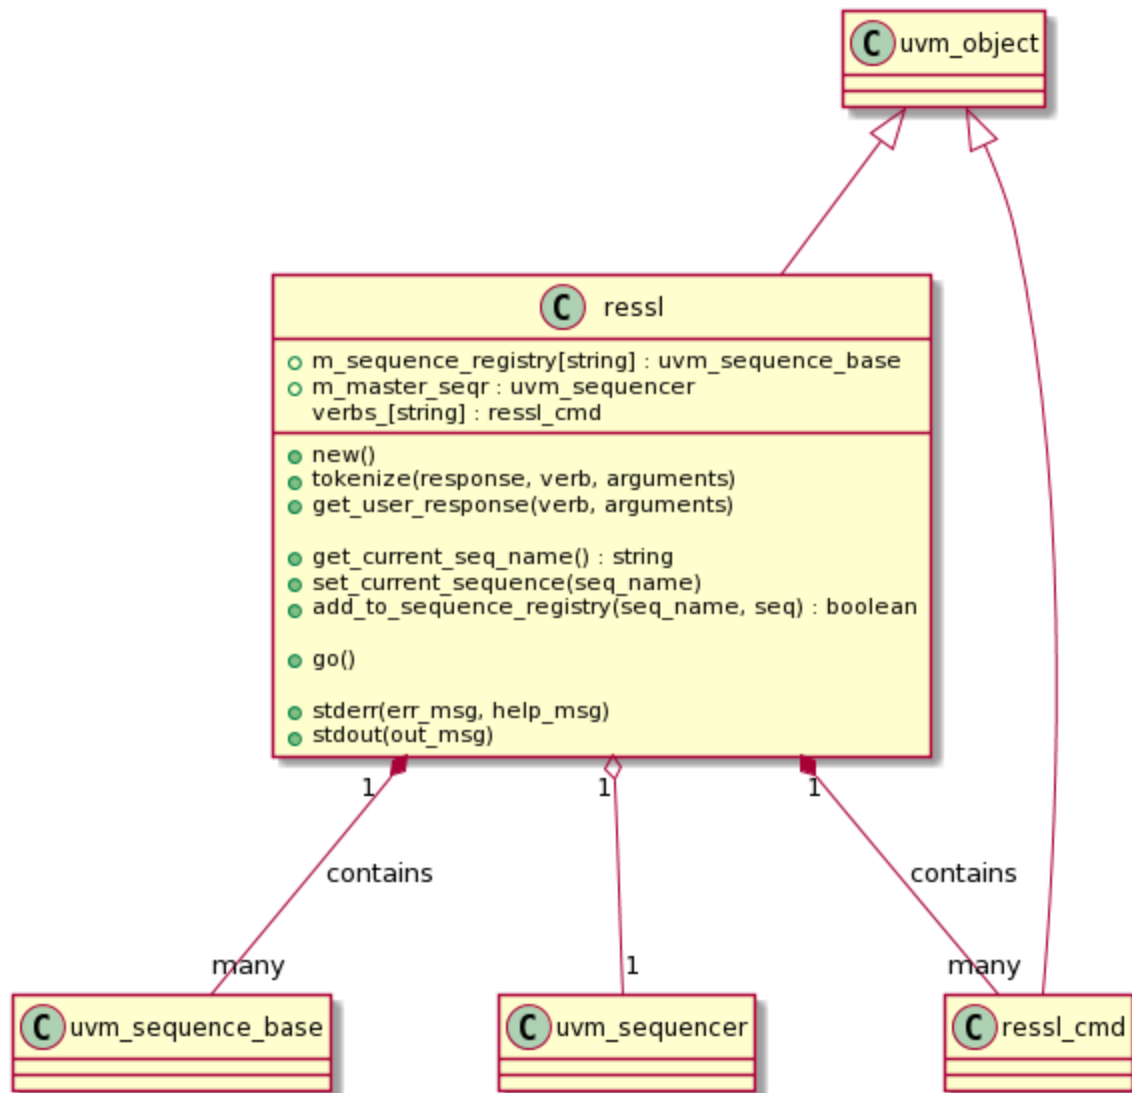


Figure 2 - RESSL Class (UML Class Diagram)

### Interpreter

The interpreter component is split into three parts:

#### Get User Input

The `ressl::go()` function implements a loop that accepts any input (one per input line) until the `quit` command is received. Each line is passed to the tokenizer for processing.

#### Tokenizer

implemented in the `ressl::tokenize()` function, it uses a simple command line parsing that splits the input into a “verb” followed by a set of “nouns”. Where the verb is assumed to be the first word in the command, and the nouns is a collection of all words that remain on the line. The `tokenize()` function accepts the user’s response and outputs the verb as a string, and the nouns as a queue of strings (can be zero). Each valid verb has an entry in the `ressl::verbs_`

associative array of handles to `ressl_cmd` objects. The queue of strings (nouns) is passed to the `ressl_cmd::execute()` command for processing. For example, the `help` command has an expected syntax of `help <cmd>`, so if the user enters `help add`, `ressl` passes the `add` string in a queue to the `help_cmd` object (derived from `ressl_cmd`) for processing.

## Verbs

Each verb in the RESSL system corresponds to a class derived from `ressl_cmd`. `ressl_cmd` implements the “strategy”[7] or policy pattern. The `ressl_cmd` is an abstract class that has two pure virtual functions:

- `execute`: performs the main processing for the command
- `help`: displays help information specific to the command.

For example, the `add` command’s `execute()` function accepts a single argument `seq_name` which indicates which sequence you want to add as a sub-sequence to the currently selected sequence. The `add` command does the following processing:

- searches the sequence registry for the `seq_name` provided and returns the handle to the associated `uvm_sequence_base`.
- using the `uvm_object::clone()` function, creates a new copy of the `uvm_sequence_base`.
- adds this cloned `uvm_sequence_base` object to the sequence registry as a sub-sequence of the currently selected sequence.

As can be seen from the UML class diagram below, there is a one-to-one correspondence between a verb in the RESSL API and a class that derives from `ressl_cmd`.

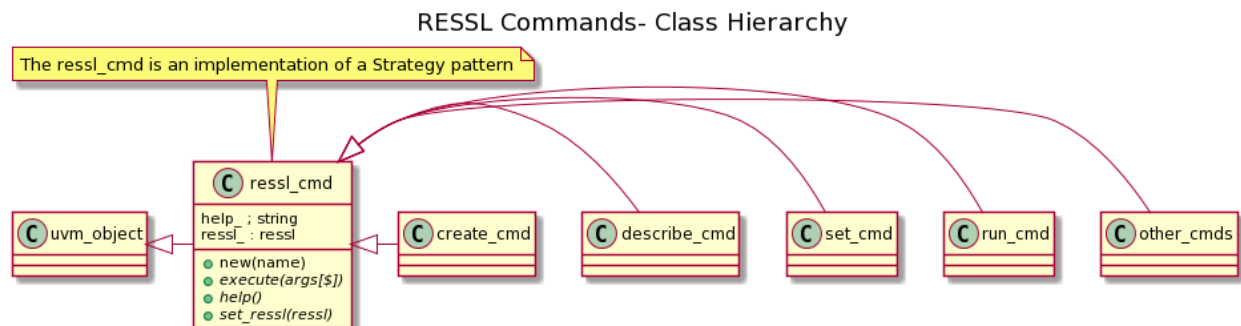


Figure 3 - RESSL Command (UML Class Diagram)

## Start Sequence

Once a new sequence has been created and added to the sequence registry it can be started using the normal `uvm_sequence::start()` function. The `ressl` object simply uses the currently selected sequence, finds it in its sequence registry and then executes a SystemVerilog `foreach` loop on each sub-sequence defined for the sequence. The UVM sequencer associated with any sequence defined in the `ressl` sequence registry must have been previously set. In the body of the `foreach` loop, this UVM sequencer is passed to the `uvm_sequence::start()` task.

## Introspection/Reflection

An important functionality for RESSL is the ability to display and set field values for any sequence in the sequence registry. RESSL implements this functionality by extending the `uvm_field` macros to add a “set” and “get” capability. This “reflection” capability required code added to the UVM 1.2 library. New functionality created in the ``uvm_field_*` macros automatically adds the ability to view and change any field defined with these macros from within RESSL.

## Storage And Retrieval

RESSL can record the sequences and sub-sequences being developed, which can then be saved to and reloaded from an external file. This enables creating and defining the sequence; provides a “replay” functionality to restore the sequence registry to the same state when the registry was “stored” to an external file. The storage file could be shared with other users, to facilitate development or to aid in debugging.




Figure 4 - RESSL External File

Once useful sequences are identified using RESSL, it is best to then code a `sequence` and add it into your verification environment (and to the instance of the RESSL). To be clear, RESSL helps identify useful scenarios by enabling a rapid development cycle that encourages exploration. At some point, the good sequences should be coded as true UVM sequences and selected during regressions.

## Interesting... but how hard is it to add to my existing environment?

To “RESSL-ize” your environment there are three key requirements, but the first two are already part of the normal development process:

- Create a set of UVM sequences that you can optionally seed the sequence registry using the `ressl::add_to_sequence_registry()` function. Generally, these are sequences of ‘atomic’ traffic, i.e., sends a stream of transactions of a certain type. These are generally

going to be the sequences that are developed as a part of normal verification, so do not represent extra work.

- Create a UVM sequencer that can be used to send these sequences. You can optionally attach these sequencers to the sequence using the `uvm_sequence_base`. This sequencer will also usually be developed as part of normal verification work.
- Add the call to the `ressl::go()` task at the appropriate time in the `run_phase` of the simulation.

To demonstrate how straightforward it is to add RESSL to your environment, we'll list the steps we did to add RESSL to the `ubus` example found in the UVM 1.2 library. Adding RESSL required changes to the base `ubus_example_base_test` and the derived `test_2m_4s` (two masters, two slaves) class.

### Step 1. Add the svlib

RESSL has been incorporated into `svlib` [9]. `svlib` is a open-source utility library for SystemVerilog that provides many useful functions including file and string manipulation, regular expression matching, etc. `svlib` uses the same open-source license as UVM and so RESSL explicitly uses the same licensing and will be freely available.

In turn, RESSL uses `svlib`'s string manipulation capabilities in the current implementation of its command tokenizer.

In order to access RESSL, we include the `svlib_macros.svh` (macro definitions used in `svlib`), and `import svlib_pkg::*`.

### Step 2. Declare and instantiate the RESSL object

Add the declaration for the `ressl` instance as part of the `ubus_example_base_test`:

```
ressl ubus_ressl;
```

Create the `ressl` instance using the UVM factory:

```
ubus_ressl = ressl::type_id::create("ubus_ressl");
```

### Step 3. Disable default sequence

The `ubus` example creates a default sequence for the masters called `loop_read_modify_write_seq` and assigns it as a default sequence for the master's sequencer. We need to "turn off" the default sequence because it interferes with the `ressl` instance's use of the master sequencer. The disabling of this default sequence is by NOT allowing the default sequence to be defined in the `uvm_config_db`. In this case, we simply comment this out, because at some point we'll likely stop using RESSL and we'd like to add this default back in (yes, there are better ways of doing this):

```
lrmw_seq = loop_read_modify_write_seq::type_id::create();  
//uvm_config_db#(uvm_sequence_base)::set(this,  
//      "ubus_example_tb0.ubus0.masters[1].sequencer.main_phase",
```

```
//          "default_sequence",
//          lrmw_seq);
```

#### Step 4. Seed RESSL's sequence registry with the `lrmw_seq`

This is an optional step as this sequence can be created using RESSL's `create` command. Nevertheless, this step populates the RESSL sequence registry with a useful sequence, and gives the user something to see when doing a `list` command at startup.

```
ubus_ressl.add_to_sequence_registry("lrmw_seq", lrmw_seq);
```

#### Step 5. Start RESSL

In the run phase of the test, we start the RESSL interpreter and pass control to it using the `ressl::go()` task:

```
ubus_ressl.go();
```

Now we're ready to start the simulation. The `go` task will block, jump out to the RESSL command line, and the user can begin working interactively as illustrated previously.

## UVM Library Modifications

In order to give RESSL users the ability to inspect and modify sequence field values, we have modified the standard UVM library. This section outlines what was added to the library, and which classes or macros are affected by these changes.

### ***Modifying `uvm_object` Class***

Specifically, `uvm_field_{int,string,enum}` macros, and the `uvm_object` class have been altered from the UVM 1.2 library (added `get_field_members` and `set_field_value` functions)

### ***Modifying UVM Field Macros***

The `uvm_field_*` macros already enable the auto-magic creation of many useful functions for each field defined including copy, compare, pack, unpack, print, and record. As discussed, in order to tweak individual elements of a sequence we needed to add the concept of introspection and reflection to RESSL. Briefly, introspection "provides the ability for a program to examine the type or properties of an object at runtime" [5]; while reflection is "the ability of a computer program to examine (see type introspection) and modify the structure and behavior of the program at runtime" [6].

For the purposes of RESSL, we need to be able to retrieve the value of any field in a sequence and display it as a string and vice-versa, accept a string and convert it to the correct type to set the corresponding field value in the sequence. For this we add conversion from and to a string as part of the UVM field macros.

Since the UVM library does not provide this capability natively, we modified the `uvm_field_*` macros. This section provides an overview of the changes to the UVM library required to add this capability.

Two new constants are added into the `uvm_object_globals` to handle the new conversion for each field:

```
parameter UVM_SETFROMSTR    = UVM_START_FUNCS+4;
parameter UVM_GETSTR        = UVM_START_FUNCS+5;
```

The `UVM_SETFROMSTR` enables the creation of method to convert *from* a string into the appropriate type defined by the field macro, e.g., `int`, `string`, `real`, etc. For `int` variables, it takes into account if a *radix* has been supplied for the field and assumes the incoming string is in the same radix, e.g., if the user sees the variable is displayed as `UVM_HEX`, the `RESSL set` command (to set the field value) is assumed to be supplied with a hex string representation, e.g., `set x 1CAFE2`.

Similarly, `UVM_GETSTR` converts the current value of the data member into a string representation. For `int` variables, it takes into account if a *radix* has been supplied and provides the correct conversion.

At this time, the introspection/reflection for `RESSL` are limited to `int`, `string`, `enum`, and `real` values. We did not add this capability to the `uvm_field_object`, and any of the array field macros.

As each field is created via the macros, it creates a new *scope* for that variable. This is contained in a composite object called `__m_uvm_status_container`. This scoping provides the ability to query which fields have been defined for the class, i.e., which scopes have been defined. So the new `UVM_SETFROMSTR` functionality searches the `__m_uvm_status_container` for the field being queried, and if found converts the string to the appropriate type value, e.g., for `int` fields:

```
field = string.atobin() for UVM_BIN
field = string.atoi()   for UVM_DEC
field = string.atohex() for UVM_HEX
```

This is done directly in the `uvm_field_int` macro found in `uvm_object_defines.svh` (macro's source).

The astute reader (which included our Verilab colleague Jonathan Bromley) will notice that the use of the `atoi<int>` functions in the conversion from `string` to `integer` has the limitation that these functions return a 32-bit `integer` value. This means that the setting of any field will be limited to 32-bit values. This limitation may be removed by the time you read this paper.

Similarly, the `UVM_GETSTR` returns the string representation, e.g., for `int` fields:

```
string = $sformatf("'b%0b", field) for UVM_BIN
string = $sformatf("'d%0d", field) for UVM_DEC
```

---

<sup>2</sup> Future update of `RESSL` will include the use of string radix “`\x###`” (hex), “`\d###`” (dec), etc. e.g., `set x \x1CAFE` to specify the radix of the new values.

```
string = $sformatf("'h%0x",field)      for UVM_HEX
```

Similar macro additions are then made for `uvm_field_string`, `uvm_field_enum` and `uvm_field_real`.

As discussed, the array based field macros and object macros do not provide this introspection/reflection capability.

## Limitations and Future Work

This section outlines both the current limitations of RESSL, and our plans going forward to mitigate these limitations and add some new functionality.

The key limitation of our RESSL system -- that can potentially prevent some teams from using RESSL to its fullest -- is that in order to get the introspection/reflection capabilities we modified the UVM 1.2 library source code (as discussed above, the `uvm_object` and the `uvm_field_*` macros were altered). The introspection/reflection capabilities are optional, as they are not required to create and execute sequences. However, this ability to see the current values for fields in sequences and be able to update them on the fly is a key feature of the system.

Otherwise, we provide a patch for the two files that are changed to add introspection/reflection. The intent is to donate this into the UVM library for subsequent releases.

### ***We do not have a way to add new constraints to a sequence.***

Any existing constraints will be respected when the sequences are randomized, but we do not support a method similar to the ``do_with()` macro. We have some thoughts on how we can provide a crude constraint mechanism, but that is not available at this time. This limitation is mitigated by the fact that the user can inspect a sequence after it has been randomized and then change values for random variables using the `set` command. Of course this means that constraints involving variables changed by the `set` command may not be satisfied.

In a practical sense, this means that the user can either set non-random variables, then randomize the sequence and then start it, or randomize the sequence, then modify variables, then start it, but not both. We do not have hooks into `pre_` or `post_randomize()` functions so any `set` commands will take place after randomization.

### ***We do not provide a way to add new fields to sequences.***

A *current* limitation is that each RESSL instance has access to exactly one sequencer. All sequences in the sequence registry must be compatible with this sequencer. We are currently testing an update to allow RESSL to hold a sequencer registry, and then the user can attach any sequence in the sequence registry to any compatible sequencer in the sequencer registry. This capability might be present in RESSL by the time we present.

## Future Features /Improvements

Future features/improvements include the following:

- redesigning the sequence registry from an associative array to a “tree” structure to allow the creation of a hierarchy of sequences. While the current system does allow nested sequences, the underlying structure to hold the sequences is an associative array -- providing only one level of hierarchy. Converting to a tree structure is both a cleaner architecture that matches the abstraction of nested sequences better, and provides opportunities for more interesting commands, e.g., copying an entire branch. This will require some additional commands or changes to existing commands to select, traverse and modify the hierarchy of sequences.
- adding mechanisms to allow constraints (even crude constraints) to be applied during randomization (a topic for a future paper, I’m sure ;-). This could include a new verb, e.g., “validate” that validates that the randomization succeeded.
- creating a “recipe” to add introspection to the sequences without altering the UVM code. This alternative requires a minimal amount of extra work while creating the sequences to ‘publicize’ the tweakable attributes.
- replacing the existing crude user command line with the VCS (Tcl) command line: by either extending the existing VCS command line with the RESSL functionality, or providing a conduit *from* RESSL to the command line. This couples the existing power of the VCS command line with the powerful capabilities of RESSL. For example, from the command line the user can force Verilog nodes during a simulation, run a sequence, evaluate the results and could interleave these activities during debug.

## Conclusions

In this paper we describe a new technique inspired by dynamic languages such as Lisp and Python to create an iterative, interactive development environment to create and start sequences on the fly called RESSL (Read-Evaluate-StartSequence-Loop). We describe the underlying design of RESSL and demonstrate that adding this environment requires little additional coding. One of RESSL’s key assets is that we have added an introspection/reflection layer to many of the UVM field macros. This allows the user to dynamically inspect and change a variable in the sequence that has been defined by a `uvm_field_*` macro. However, this introspection/reflection capability is also a key limitation for many users since it requires you to use a modified UVM library.

We also describe some typical user scenarios, existing limitations of the system, and identify some potential future work.

Using the RESSL system enables an interactive development environment that encourages the easy and iterative development of a sequence library without having to recompile the code. Using RESSL provides a platform where RTL designers are given a verification environment where they can drive simple sequences, easily change sequence variables, and create more complex nested sequences to test their design. For verification engineers it allows them to interactively explore and define interesting sequences in an iterative way, and create libraries of interesting scenarios.



RESSL will continue to evolve and develop as we learn more about using such an interactive, dynamic way to develop sequences. We've identified some future improvements such as allowing sequencers to be created and attached to sequences, permitting constraints to be applied, etc. RESSL is freely available with the same open-source licensing as UVM and we would welcome comments, ideas, and patches for inclusion into future revisions.

All in all, using a system like RESSL gives you the ability to tag-team in a cage-match, launch UVM sequences off the turnbuckle and wrestle the DUT into a submission hold with no hope for re-match to help you become known as “The Bug *Crusher*”. Stay tuned for more “Smack-Down” action as we continue developing RESSL.

## Acknowledgements

We would like to thank our colleagues at Verilab for their careful and positive review comments. With special thanks to Jonathan Bromley for his development of the svlib, his encouragement of this idea, and the generous gift of his time for the questions that popped up while writing this paper.

## References

- [1] Python Language Website: <https://www.python.org/>
- [2] Ruby Language Website: <https://www.ruby-lang.org/en/>
- [3] Wikipedia entry for Lisp: [http://en.wikipedia.org/wiki/Lisp\\_%28programming\\_language%29](http://en.wikipedia.org/wiki/Lisp_%28programming_language%29)
- [4] Wikipedia entry describing the software Factory Pattern: [http://en.wikipedia.org/wiki/Factory\\_method\\_pattern](http://en.wikipedia.org/wiki/Factory_method_pattern)
- [5] Wikipedia entry describing type introspection: [http://en.wikipedia.org/wiki/Type\\_introspection](http://en.wikipedia.org/wiki/Type_introspection)
- [6] Wikipedia entry describing Reflection: [http://en.wikipedia.org/wiki/Reflection\\_\(computer\\_programming\)](http://en.wikipedia.org/wiki/Reflection_(computer_programming))
- [7] Wikipedia entry described Strategy software pattern: [http://en.wikipedia.org/wiki/Strategy\\_pattern](http://en.wikipedia.org/wiki/Strategy_pattern)
- [8] Lisp I Programmer's Manual (1960), References to REPL on page 2, [http://history.siam.org/sup/Fox\\_1960\\_LISP.pdf](http://history.siam.org/sup/Fox_1960_LISP.pdf)
- [9] svlib - a programmer's utility library for SystemVerilog written by Jonathan Bromley, <http://www.verilab.com/resources/svlib/>